

Flux Quiz by CCL

1. Which of the following is NOT true about the MNIST dataset in the lecture? (time stamp: 1:24)
 - ☐ Handwritten digits were extracted from cheques.
 - ☐ Can be loaded in Flux via `Flux.Data.MNIST.labels()/images()`.
 - ☐ Consists of RGB images of dimensions (28,28,3)
 - ☐ Contains a train:test split of 5000:55000 images.
2. Which of the following is NOT a function of defining a test set of 5000 images? (time stamp: 5:10)
 - ☐ The relatively small testset can reduce the training time of each epoch.
 - ☐ The testset helps estimate the model's generalisation capabilities.
 - ☐ The testset tests for signs of model overfitting (test accuracy \ll train accuracy)
 - ☐ The testset will be integrated into the trainset at later epochs for finetuning.
3. What is the output of the softmax function in `model = Chain(Dense(n_inputs, n_outputs, identity), softmax)`? (time stamp: 7:37)
 - ☐ 1 discrete label from the 0-9.
 - ☐ A 10 by 10 feature map.
 - ☐ 10 probabilities representing the likelihood that the input corresponds to each class.
 - ☐ A `10,1` one-hot vector, where the 1-value corresponds to the predicted class.
4. Which of the following is NOT true about the cross entropy (CE) function in `L(x,y)=Flux.crossentropy(model(x), y)`? (time stamp: 7:53)
 - ☐ It calculates calculate the loss of the model.
 - ☐ It compares probability distributions.
 - ☐ It is equivalent to the log loss function.
 - ☐ It is a convex function.
5. What do a,b,c refer to respectively in `Flux.train!(a, [b], c)`? (time stamp: 8:52)
 - ☐ CE loss — training batch — SGD optimiser
 - ☐ training set — CE loss — SGD optimiser
 - ☐ training set — CE loss — Adam optimiser
 - ☐ Nesterov optimiser — CE loss — training batch
6. How is the throttle function applied to Flux callbacks? (time stamp: 10:30)
 - ☐ To pause the model training.
 - ☐ To update the learning rate automatically.
 - ☐ To redirect Jupyter Notebook output from stdout to a log file.
 - ☐ To control how frequently the callback function is executed.
7. Which is NOT a possible cause of having a test loss lower than a train loss? (time stamp: 13:55)
 - ☐ An overfitted model.
 - ☐ Test data being accidentally included in train set.
 - ☐ An incorrectly written loss function.
 - ☐ An underfitted model.

8. What does the number 20 represent in `model = Chain(Dense(n_inputs, 20, relu), Dense(20, n_outputs, identity), softmax)`? (time stamp: 18:55)
- ☐ The total number of possible output classes.
 - ☐ The number of neurons in the hidden layer.
 - ☐ The number of hidden layers.
 - ☐ The spatial dimensions of input MNIST images.
9. Which of the following is NOT true about the ReLU function? (time stamp: 19:23)
- ☐ It is a step function.
 - ☐ It is non-linear.
 - ☐ It increases model complexity.
 - ☐ It is a saturating function.
10. Which of the following is NOT a way to counter overfitting problems? (time stamp: 21:33)
- ☐ Use more training data.
 - ☐ Use more test data.
 - ☐ Reduce the number of model parameters.
 - ☐ Shuffling data before defining train and test sets.

Answers: C, D, C, C, A, D, A, B, D, B