

| |
|-------------------------|
| Knet Quiz by CCL |
|-------------------------|

1. Lecture 2: Which of the following is NOT true about Visual Question Answering (VQA)? (time stamp: 4:10)
 - ☐ Models take both sequential (text) and discrete (image) data as input.
 - ☐ Models use attention mechanisms to heed corresponding image regions and text.
 - ☐ Models are trained recurrently via time step iterations.
 - ☐ Models pay equal attention to every word of the input sentence to avoid biases.
2. Lecture 2: Which of the following is NOT true about attention mechanisms in a word-level VQA model? (time stamp: 4:10)
 - ☐ Allows humans to better understand AI's reasoning process and inference results.
 - ☐ Contains information on the corresponding image regions of individual words.
 - ☐ The relationship between corresponding image regions and words is a one-to-one mapping.
 - ☐ Attention can be visualised through plotting attention maps.
3. Lecture 3: Which of the following is NOT true about Knet's built-in MNIST dataset? (time stamp: 2:29)
 - ☐ Uses mini-batches of 100.
 - ☐ Uses a channels-first image representation.
 - ☐ Spatial dimensions (height, width) of images are 28x28.
 - ☐ There are in total 10 possible predicted classes.
4. Lecture 3: In `y = f.w * mat(x) .+ f.b`, if x is of dimensions (28,28,1) and y is (64,1), what are the dimensions of f.w, f.b respectively? (time stamp: 3:35)
 - ☐ (64,28,28) — (64,28,28)
 - ☐ (64,28,28) — (64,1)
 - ☐ (64,784) — (64,1)
 - ☐ (64,784) — (64,28,28)
5. Lecture 3: Which of the following is a correct description of the LeNet architecture used? (time stamp: 3:46)
 - ☐ `lenet = RNN((784, 1), 256, 64, 10)`
 - ☐ `lenet = CNN((28,28,1), (14,14,16), (7,7,32), 448, 10)`
 - ☐ `lenet = CNN((28,28,1), (5,5,20), (5,5,20), 500, 10)`
 - ☐ `lenet = CNN((28,28,3), (5,5,20), (5,5,20), 25, 10)`
6. Lecture 4: What do `<unk>` tokens mainly represent? (time stamp: 2:22)
 - ☐ Words with characters which cannot be encoded in UTF-8.
 - ☐ Rarely seen words which are not indexed in our corpus.
 - ☐ Words of another language, e.g. Chinese characters.
 - ☐ Expletives which are sanitised from our dataset.
7. Lecture 4: What is NOT a function of the test set? (time stamp: 2:47)
 - ☐ To measure the accuracy of our model.
 - ☐ To check for model overfitting.

- ☐ To be integrate into the trainset at later iterations for model finetuning.
 - ☐ To assess specific features of the model, such as false positives, false negatives and recall.
8. Lecture 5: Which of the following is NOT true about character-level RNNs? (time stamp: 0:47)
- ☐ 1 new character is predicted at each time step.
 - ☐ The n'th prediction is generated using the (n-1)'th prediction, previous & current hidden states.
 - ☐ A "one to one" character-level RNN can be used to generate one paragraph of text.
 - ☐ Input characters are represented by one-hot vectors.
9. Lecture 5: Which of the following is NOT an insight learned by character-level RNNs? (time stamp: 3:12)
- ☐ The semantic meaning of words and sentences.
 - ☐ Capitalisation rules in English syntax.
 - ☐ Play transcript formats where character names are put above their lines.
 - ☐ Basic punctuation of sentences.
10. Lecture 6: Which of the following is NOT information given by examining a model's recall@5 (top-5 classification error)? (time stamp: 3:36)
- ☐ To determine the intraclass variance inferred by the model.
 - ☐ To determine the interclass variance inferred by the model.
 - ☐ To assess performance on specific data classes.
 - ☐ To assess the ratio of true positive data, to false positives and false negatives.

Answers: D, C, B, C, C, B, C, C, A, A